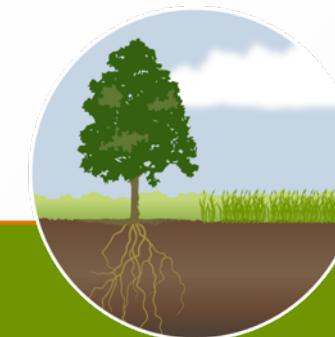
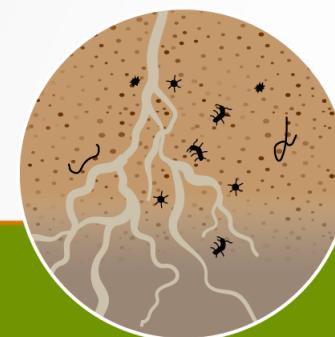


The Road to DevOps HPC Cluster Management

Ken Schmidt, Evan Felix

Nov 22, 2019



SCIENTIFIC INNOVATION THROUGH INTEGRATION

Short EMSL Computing History



- MPP2 – First Generation Cluster
 - ▶ One of the first Linux Clusters in the top 5
 - ▶ One of the earliest deployments Lustre
 - ▶ ~1000 nodes – How do you install it?
 - ▶ More than a year for deployment and acceptance
- Chinook – Second Generation
 - ▶ Let's do it again – only bigger!
 - ▶ Deployment took months
 - ▶ Time to strike out on our own
 - Drop vendor product
 - Use open source tools – xCAT





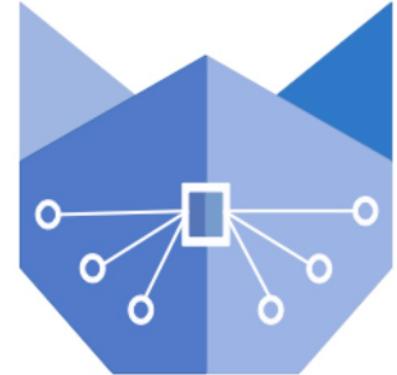
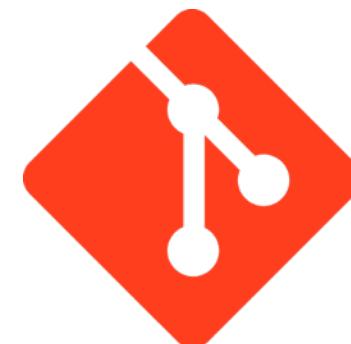
- Cascade – Third Generation
 - ▶ Hardware only procurement
 - ▶ Same configuration repository
 - ▶ Deployed with Scientific Linux 6
 - ▶ From delivery to acceptance to full production in weeks
 - ▶ Upgraded to CentOS 7 without user problems





Open Source advantages

- xCAT
 - ▶ Source is available to modify and contribute back changes
 - ▶ Decouple hardware, operating system, and management software
- Git
 - ▶ Track all changes made to xCAT database
 - ▶ Store all configurations and scripts
 - ▶ Facilitate collaboration between administrators
- Gitlab
 - ▶ Facilitate collaboration between administrators
 - ▶ Track work being done
 - ▶ Store Documentation for new/unfamiliar administrators



Monthly Outages



- Every month take an outage
 - ▶ Makes the system more up-to-date for the users
 - ▶ Reduces the number of unplanned outages
 - ▶ Gives administrators full unfettered access to make changes
 - ▶ Drives development cycles
 - ▶ ~1.5% yearly downtime

Weekly Meetings



- Week 1
 - ▶ What work was completed during the last outage? – outage minutes
 - ▶ What went wrong that needs to be changed?
 - ▶ What can be done better?
- Week 2
 - ▶ Plan the tasks/priorities for up upcoming outage
- Week 3
 - ▶ Evaluate what will be complete by outage
 - ▶ Schedule outage (plan resources & check for interdependencies)
- Week 4
 - ▶ Burn the cluster to the ground and build it up again



- Issue Tracking
 - ▶ All work is tracked in Gitlab using issues
 - ▶ Progress is tracked to determine what will and what won't go into the upcoming outage
 - ▶ Admin/Developer only use
- Documentation
 - ▶ All procedures are documented
 - ▶ They are never complete and need almost constant updating
 - ▶ Gitlab provides easy web and markdown file wikis
 - ▶ Procedures are occasionally tested by outside driver outages

Disadvantages of Traditional Management



- “Organically” Grown – Install a system and make changes as needed
 - ▶ Upgrades are near impossible
 - ▶ “How did we get here?”
 - ▶ “How do we upgrade?”
 - ▶ “How do we prevent user interruption?”



Working with (around) xCAT

- Automated table backups checked into git
- Difficult to test changes on test cluster and merge into production
- Move configurations out of xCAT's databases into revision control
- Use glue script to jump from xCAT config into scripts stored in revision control

```
#!/bin/sh

set -x # Log everything the script is doing
scriptdir=$1
shift

for i in $scriptdir/* # For every file in the directory passed on the command line.
do
    if [ -x $i ] # Make sure the script is executable
    then
        if ! $i $* # Execute postscript and check the results
        then
            set +x                      # If the script fails,
            echo "Error executing $i postscript" # log the error
            exit 1                         # and bail
            set -x
        fi
    done
fi
```

```
.
.
00_savenodeinfo -> /install/postscripts/msc/custom/savenodeinfo
01_confignics -> /install/postscripts/confignics
05_mkresolvconf -> /install/postscripts/mkresolvconf
10_syncfiles -> /install/postscripts/syncfiles
15_configxcathost -> /install/postscripts/configxcathost
20_setup_sssd -> /install/postscripts/msc/custom/setup_sssd
30_hardeths -> /install/postscripts/hardeths
40_rsyslog -> /install/postscripts/msc/custom/rsyslog
50_remoteshell -> /install/postscripts/msc/custom/remoteshell
60_mountscratch -> /install/postscripts/msc/custom/mountscratch
70_rc-local -> /install/postscripts/msc/custom/rc-local
80_hugepages -> /install/postscripts/hugepages
```

MSC's Implementation of Cluster as Code



- All configurations are checked into repositories and are checked out minutely on the clusters' management nodes
 - ▶ Any change attempted outside git are stashed and removed
- Any change gets committed to the development branch and is tested on test cluster
 - ▶ Prior to outages, development is merged into production
- Easily track changes to configurations and can track down when/why the change was made
- Example on github - <https://github.com/EMSL-MSC/MscClusterMgmt>



- New Cluster – Next Generation
 - ▶ Hardware only procurement
 - ▶ Will be deployed with CentOS 7
 - ▶ Planned upgrade to CentOS 8 in its lifetime
 - ▶ Same configuration repository
 - ▶ Expected smooth transition based on past performance

Questions?

The research was performed using EMSL (grid.436923.9), a DOE Office of Science User Facility sponsored by the Office of Biological and Environmental Research.



kenneth.schmidt@pnnl.gov
www.emsl.pnnl.gov

